

CS556 - ITERATIVE AND MULTIGRID METHODS: FALL 2018  
ALGEBRAIC MULTIGRID METHODS FOR NONLINEAR  
HYPERELASTICITY PROBLEMS

BHAVESH SHRIMALI TEJASWIN PARTHASARATHY \*

**Abstract.** We survey the performance of algebraic multigrid (AMG) methods to solve nonlinear elliptic partial differential equations (PDEs) in hyperelasticity. We discretize the underlying PDE in  $N = 2$  space dimensions using conforming quadrilateral ( $Q-k$ ) finite elements and solve the resulting algebraic equations for the nodal displacements using two different AMG algorithms (a): Newton-AMG and (b) Full Approximation Scheme (FAS). We explore the effectiveness of AMG across a breadth of physical (e.g. material compressibility) and algorithmic (e.g. choice of near null-space candidate vectors) parameters. We highlight the potential benefits of employing AMG, both as a solver and preconditioner, for hyperelastic problems.

**1. Problem Statement.** The effective finite element solution of elliptic PDEs arising in solid mechanics, e.g. hyperelasticity (10), homogenization (8) and electro-magneto-elasticity (9) requires highly efficient and parallelizable iterative solvers. In this paper we seek to develop numerical solutions to a sample 2D BVP in hyperelasticity with the following choice of free energy function  $W(\mathbf{F})$

$$(1.1) \quad W(\mathbf{F}) := \frac{\mu}{2} (I_1 - 2) - \mu \ln J + \frac{\lambda}{2} (J - 1)^2, \quad \text{where} \quad I_1 = \mathbf{F} \cdot \mathbf{F}, \quad J = \det \mathbf{F}.$$

20 The corresponding constitutive relation and the strong form of the governing equation  
 21 can then be compactly written as

$$\begin{aligned}
 & (1.2) \quad \mathbf{S}(\mathbf{F}; \mathbf{X}) = \frac{\partial W}{\partial \mathbf{F}}(\mathbf{F}; \mathbf{X}) = \mu \mathbf{F} - \mu \mathbf{F}^{-T} + \lambda J(J-1) \mathbf{F}^{-T} \\
 & (1.3) \quad \nabla \mathbf{S} + \mathbf{b}_0 = \mathbf{0} \quad \text{in } \Omega \quad \text{and} \quad \begin{cases} \chi(\mathbf{X}) = \bar{\mathbf{F}}\mathbf{X}, & \text{for } \mathbf{X} \in \partial\Omega_x, \\ \mathbf{S}\mathbf{n}_0 = \mathbf{s}_0 & \text{for } \mathbf{X} \in \partial\Omega_t, \end{cases}
 \end{aligned}$$

where  $\mathbf{X}$  represents the Lagrangian coordinate,  $\mathbf{F} = \nabla\chi(\mathbf{X})$  is called the deformation gradient,  $\mathbf{S}$  the first Piola-Kirchoff stress tensor,  $\mu$  the initial shear modulus, and  $\lambda$  the lamé constant. The body force  $\mathbf{b}_0$  in (1.3) is an important parameter to develop and test benchmark solutions. Alternatively the problem can also be expressed in the form of a variational principle

$$(1.4) \quad \min_{u \in \mathcal{U}} \Pi(u) := \int_{\Omega} W(\mathbf{F}) d\mathbf{X} - \int_{\Omega} \mathbf{b}_0 \cdot \mathbf{u} d\mathbf{X} - \int_{\partial\Omega_s} \mathbf{s}_0 \cdot \mathbf{u} d\mathbf{X}.$$

32 where

$$(1.5) \quad \mathcal{U} = \mathbf{v} \ni \mathbf{v} \in W^{(1,k)}(\Omega), \mathbf{v} = \bar{\mathbf{F}}\mathbf{X} \quad \forall \mathbf{X} \in \partial\Omega_x$$

35 is a set of candidate solutions satisfying kinematic boundary conditions. The Euler-  
 36 Lagrange equations associated with the above variational principle are given by

$$(1.6) \quad \mathcal{G}(\mathbf{u}, \mathbf{v}) = \int_{\Omega} \frac{\partial W}{\partial \mathbf{F}}(\nabla \mathbf{u}; \mathbf{X}) \cdot \delta \nabla \mathbf{v} \, d\mathbf{X} - \int_{\Omega} \mathbf{b}_0 \cdot \delta \mathbf{v} \, d\mathbf{X} - \int_{\partial \Omega_t} \mathbf{s}_0 \cdot \delta \mathbf{v} \, d\mathbf{X} = 0$$

\*bshrima2@illinois.edu (Bhavesh Shrimali)    tp5@illinois.edu (Tejaswin Parthasarathy)

40 where

41 (1.7)  $\mathcal{U}_0 = \mathbf{v} \ni \mathbf{v} \in W^{(1,k)}(\Omega), \mathbf{v} = \mathbf{0} \quad \forall \mathbf{X} \in \partial\Omega_x$

43 Choosing a suitable partition of the domain  $\Omega$  into  $\mathcal{N}$  non-overlapping elements and  
44 a polynomial basis  $\varphi(\mathbf{X})$  with compact support over each element to represent the  
45 primal field  $\mathbf{u} = \varphi(\mathbf{X}) \mathbf{d}$  and the test function  $\mathbf{v} = \varphi(\mathbf{X}) \mathbf{d}_0$  respectively, the discrete  
46 version of (1.6) and its linearized form (using Newton's method) with a slight abuse  
47 of notation, can be written as

48 (1.8)  $\mathcal{G}(\mathbf{d}, \mathbf{d}_0) = 0 \quad \forall \mathbf{d}_0$

49 (1.9)  $\mathcal{G}(\mathbf{d}^k, \mathbf{d}_0) + D\mathcal{G}(\mathbf{d}^k, \mathbf{d}_0) \delta\mathbf{d}^k = 0 \quad \forall \mathbf{d}_0$

51 **2. Approach.** We solve the resulting algebraic system (1.8) using non-linear  
52 AMG methods. Our motivation to use these techniques stems from their efficiency (3),  
53 high parallelizability (1; 2) and a spate of successful applications in solving such non-  
54 linear problems (4; 6), with good convergence properties and small time-to-solution  
55 (TTS). Moreover the discretized systems considered in this work are expected to be  
56 SPD, based on the energy minimization property (1.4), for which AMG methods have  
57 a history of success (3; 5). We focus only on two such methods—(a) Newton–AMG  
58 and (b) Full Approximation Scheme (FAS). Newton–AMG involves approximating  
59 the solution to (1.8), by equivalently solving a series of locally-valid linear systems  
60 (1.9) using AMG. FAS however, solves (1.8) by solving a series of smaller non-linear  
61 problems, using concepts of AMG.

62 For both the methods listed above, we utilize the method of manufactured solution  
63 to test the validity of our solver(s). We proceed to study the effect of physical  
64 (for e.g.  $\lambda$ , the lamé constant representing material compressibility) and algorithmic  
65 parameters (for e.g. the choice of near-null space candidate vectors) on the perfor-  
66 mance of AMG as both a solver and a pre-conditioner, which include the convergence  
67 factors and TTS. We then attempt to explain our observations based on physical  
68 interpretations.

69 We will perform the computation in Python, using vectorized code from `numpy`  
70 for dense- and `scipy.sparse` for sparse-matrix/vector operations, `sympy` for symbolic  
71 algebraic manipulations and `matplotlib`, VTK for visualization. We also utilize the  
72 AMG routines and canned solvers in PyAMG (7).

73 **3. Implementation and Numerical Results.** In this section we briefly de-  
74 scribe the physical and numerical (conforming FE) framework and present repre-  
75 sentative simulations. We take a prototypical 2D domain in the Euclidean space  
76  $\Omega \subset \mathbb{R}^2 = [0, 10]^2$  for solving (1.3) and subject it to various prescribed boundary  
77 conditions of the form (1.3)<sub>2</sub> by changing  $\bar{\mathbf{F}}$ . The resulting geometric nonlinearities  
78 are successfully captured by our solver (see Figure 3.1) in the sense that it generates  
79 non-singular algebraic systems. Additionally we construct the algorithm such that it  
80 can be used for different variational problems (1.4) by directly changing the free- or  
81 stored-energy (1.1) in near-mathematical notation, without tweaking other parts of  
82 the algorithmic pipeline.

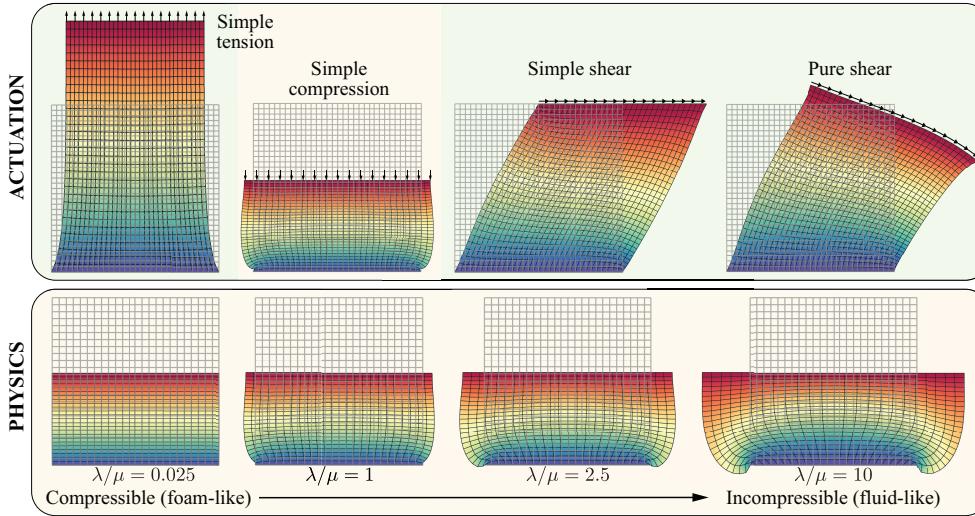


Fig. 3.1: Solver capability: Our FE strategy works across a range of actuation (prescribed displacements with stretches as high as 150%) and physical (prescribed as changes in the constitutive model) parameters. The colors represent the material deformation in the direction of these prescribed displacements (indicated by arrows on the corresponding boundary), with blue and red representing low and high deformations respectively.

We then demonstrate the capabilities of our FE solver to reproduce manufactured solutions of (1.3). The physical setup that reproduces one such solution, with  $\bar{\mathbf{F}} = \text{diag} [\lambda_1, \bar{\lambda} = 1.5]$  is shown in Figure 3.2(a) and the corresponding non-linear deformation, captured using  $Q-1$  elements, is shown in Figure 3.1(b). With this setting we observe a first-order spatial convergence, as seen in Figure 3.1(c). We add that the deformation is prescribed quasi-statically in steps (see inset, Figure 3.1(d)), with the nonlinear system of equations (1.8) at each step solved using Newton's method, which exhibits quadratic convergence (Figure 3.1(d)).

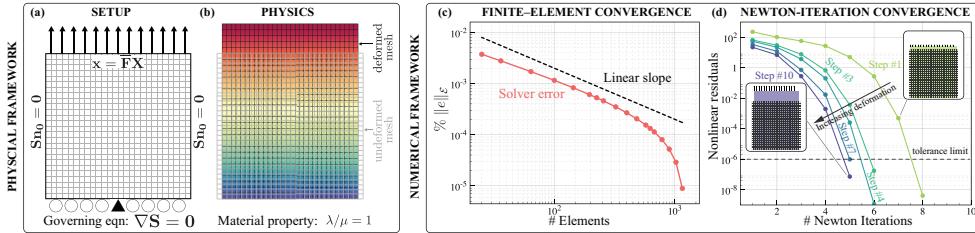


Fig. 3.2: Solver validity: We validate our solution for the simple tension test case shown in (a) against its manufactured solution, wherein we solve  $\nabla \mathbf{S} = \mathbf{0}$  with prescribed displacement at the top (indicated by arrows). This leads to (b) high deformations, with red (blue) indicating high (low) values in the imposed direction. (c) By tracking the relative error in the energy norm  $\|\cdot\|_{\mathcal{E}}$  (see ?? for details) with increasing spatial resolution, we observe that FEM converges linearly. To ensure global convergence, we prescribe deformations distributed among many steps (10 in this case), leading to (d, inset) quasi-static material deformations. We use Newton's method to perform the non-linear solve at each such step, and observe the (d) quadratic convergence at *every* step as expected.

**3.1. Newton–AMG.** We solve the linearized system (1.9) using AMG. We choose a model pure-shear problem (see Figure 3.1) in view of its (a) simplicity and (b) ability to incorporate compression, extension and shearing of elements from

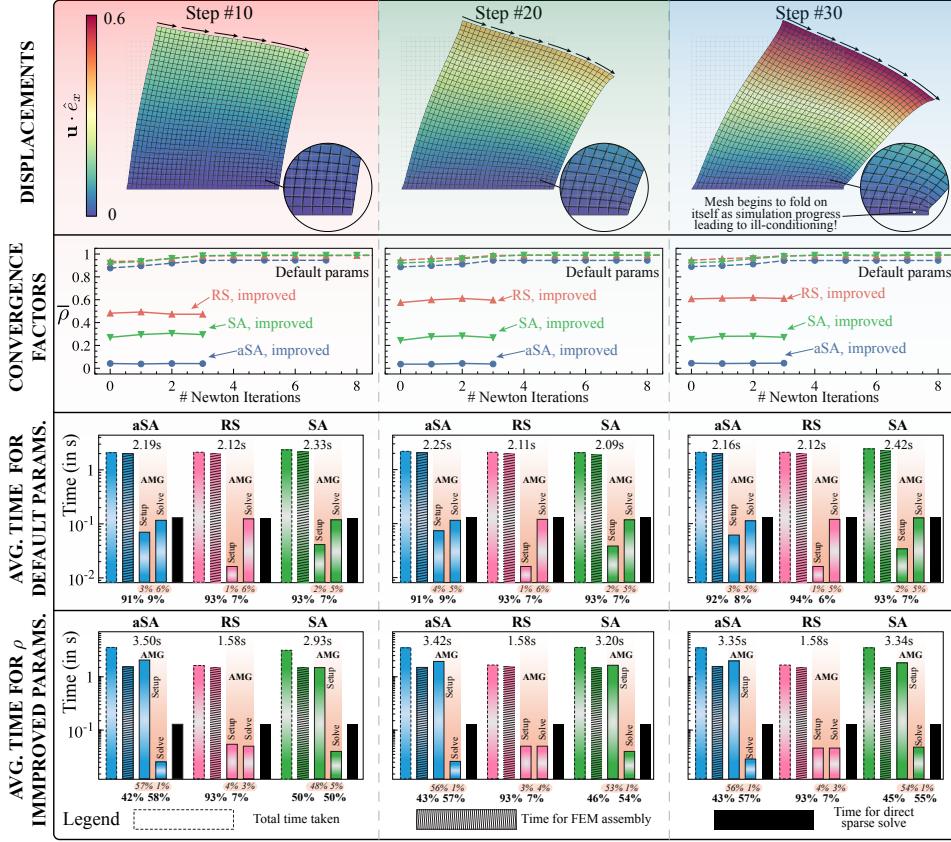


Fig. 3.3: AMG: To compare the performance of AMG as a preconditioner to CG in Newton’s method, we simulate the pure-shear case shown in (row 1) above. Here we go from a well-conditioned system (in step #10, left) to an ill-conditioned system (in step #30, right), due to mesh folding. The colors indicate the x-displacement of the material, from blue (no displacement) to red (max displacement of 0.6). Using AMG, we see (row 2) bad performance in the convergence factors, using the default parameters of PyAMG, indicated by the dashed lines. We iterate on the parameter set to arrive at an improved set, with better performance (solid lines) for Ruge–Stüben (red lines), Smoothed Aggregation (SA, green lines) and adaptive Smoothed Aggregation (aSA, blue lines) preconditioners. We instrument the code and obtain average timing data for both default (row 3) and improved (row 4) parameters, represented as a bar plot in the semilog scale. Here dashed borders for the bar indicate total time, a hatched bar indicates assembly time and black bars represent the time taken to solve by the direct sparse solver routine of `Scipy.sparse`. The AMG times are marked with an orange halo, with both setup and solve times listed. The percentage split up between different components (and within each component, whenever possible is also shown).

94 characteristic physics, thus testing all ‘modes’, while simultaneously (c) enabling the  
 95 comparison of AMG performance across a range of linear-system conditioning, from  
 96 well-conditioned (step #10, #20 in the first row of Figure 3.3) to ill-conditioned (step  
 97 #30 in the first row of Figure 3.3) systems due to eventual folding of the mesh with  
 98 increasing deformation.

99 We obtain the solutions of the linear system using preconditioned conjugate-  
 100 gradient with three different AMG preconditioning algorithms—(a) adaptive Smoothed  
 101 Aggregation (aSA), (b) Ruge–Stüben with classical C/F splitting and (c) Smoothed  
 102 Aggregation (SA)—and compare their performance metrics at the three afore-  
 103 mentioned load-steps. The first metric relates to the averaged convergence factors  $\bar{\rho}$  (de-

104 fined in the [Algorithm A.1](#)) for every Newton iteration and is shown in the second row  
105 of [Figure 3.3](#). With the default parameter set from PyAMG, the  $\bar{\rho}$  for all three variants  
106  $\in [0.9, 0.95]$ , indicating an ineffective solver. We note that in some cases an increase  
107 in Newton Iterations is quintessential in reducing the residuals (till convergence) of  
108 the non-linear equations [\(1.8\)](#), as the maximum number of Krylov iterations within  
109 one Newton iteration was restricted.

110 We then seek to improve  $\bar{\rho}$  using different set of parameters for each of the above  
111 algorithms. We report them in the same graph, with the corresponding improved  
112 parameter set reported in [Table A.1](#). Not surprisingly, we observe improvement in  
113 performance, with aSA leading in the  $\bar{\rho}$  metric, followed by SA and RS. We further  
114 observe remarkable consistency in convergence across conditioning (from step#10 to  
115 #30), with only a mild penalty for the ill-conditioned cases. Moreover, aSA has text-  
116 book performance ( $\bar{\rho} = 0.04 - 0.06$ ), presumably because it finds the best null-space  
117 candidates, corresponding to the rigid translation and rotation modes of the sys-  
118 tem. SA, with manually estimated translation–rotation null-space candidates, suffers  
119 a penalty of  $5 \times$  in  $\bar{\rho}$ .

120 Investigating the time taken for obtaining the solutions for *one load step*, we  
121 obtain the third row of [Figure 3.3](#), with the total, assembly, setup and solve times  
122 (and their percentage in the composition) reported for different AMG solvers for the  
123 default parameters. Across the different load steps and solvers, we notice that the total  
124 time taken is similar, with the bottleneck of our algorithmic implementation lying in  
125 the matrix assembly—indeed, we loop over all elements in our Python implementation,  
126 leading to characteristic computational inefficiencies. The RS preconditioner takes the  
127 least amount of time to setup, followed by SA (due to aggregation) and finally aSA  
128 (to adapt, aka form the candidate vectors). The solve times, however, are similar. A  
129 comparison with the direct sparse solver from `Scipy.sparse` indicates that the AMG  
130 solve is slightly more efficient, but comes at the expense of the required setup time.

131 The parameter set for the cases with improved convergence changes the time split-  
132 up. We now see (fourth row of [Figure 3.3](#)) that aSA takes the most amount of time,  
133 followed by SA and then RS, in contrast with similar total times observed earlier in  
134 row three. We directly correlate this to the setup times—the need for more adaptation  
135 (more candidate vectors), smoothing the interpolation operators etc. increase heavily  
136 the setup cost of aSA and SA. RS also sees an increased setup time, but not to the  
137 same extent as its counterparts. The cost incurred in the setup pays off during the  
138 solve, with exactly the reverse order—RS is more expensive, followed by SA and then  
139 aSA. We add that all AMG routines perform far better than a direct solver, once  
140 setup. An interesting observation pertains to the TTS—while aSA performs better in  
141 the  $\bar{\rho}$  metric, RS performs the best in the metric of total time taken (which is of the  
142 most interest), taking similar time (in both solve and setup) to a direct sparse solve,  
143 for the system size investigated.

144 **4. Conclusions.** In this paper we investigated the performance of AMG precon-  
145 conditioners for solving system of nonlinear algebraic equations encountered in Hyper-  
146 elasticity. Newton–AMG was chosen as the method to solve the nonlinear equations  
147 in view of (a) robustness of Newton’s method to deliver quadratic convergence for  
148 variational problems, and (b) ability of AMG to provide accurate preconditioners  
149 to iterative Krylov-subspace solvers (`cg`, `gmres`, etc.). We recorded the performance  
150 of adaptive Smoothed Aggregation, Classical Ruge-Stüben and Smoothed Aggrega-  
151 tion methods in solving a representative BVP. Our  $\bar{\rho}$  experiments indicated that aSA  
152 can have remarkable performance—this feeds into the idea of having *AMG-memory*,

153 wherein we take the adaptation information from a prior iteration and feed it into  
 154 consequent iteration(s), to potentially reduce setup times. From our timing experi-  
 155 ments, we infer that balancing setup and solve times of AMG might be key to obtain  
 156 improved performance—indeed this was the case with our improved RS parameter set,  
 157 which had better TTS while still converging within four Newton iterations. While  
 158 these methods perform reasonably well when compared to the available direct solvers,  
 159 their full potential can be achieved for larger problem sizes and with a more robust  
 160 framework.

## 161 References.

- 162 [1] M. F. ADAMS, H. H. BAYRAKTAR, T. M. KEAVENY, AND P. PAPADOPOULOS,  
*163 Ultrascalable implicit finite element analyses in solid mechanics with over a half  
 164 a billion degrees of freedom*, in Proceedings of the 2004 ACM/IEEE conference  
 165 on Supercomputing, IEEE Computer Society, 2004, p. 34.
- 166 [2] A. H. BAKER, R. D. FALGOUT, T. V. KOLEV, AND U. M. YANG, *Scaling  
 167 hypre’s multigrid solvers to 100,000 cores*, in High-Performance Scientific Com-  
 168 puting, Springer, 2012, pp. 261–279.
- 169 [3] W. L. BRIGGS, S. F. MCCORMICK, ET AL., *A multigrid tutorial*, vol. 72, Siam,  
 170 2000.
- 171 [4] M. W. GEE AND R. S. TUMINARO, *Nonlinear algebraic multigrid for constrained  
 172 solid mechanics problems using trilinos.*, tech. report, Sandia National Labora-  
 173 tories, 2006.
- 174 [5] V. E. HENSON, *Multigrid methods for nonlinear problems: an overview*, tech.  
 175 report, Lawrence Livermore National Lab., CA (US), 2002.
- 176 [6] A. KLAWONN, M. LANSER, AND O. RHEINBACH, *Toward extremely scalable  
 177 nonlinear domain decomposition methods for elliptic partial differential equations*,  
 178 SIAM Journal on Scientific Computing, 37 (2015), pp. C667–C696.
- 179 [7] L. N. OLSON AND J. B. SCHRODER, *PyAMG: Algebraic multigrid solvers in  
 180 Python v4.0*, 2018, <https://github.com/pyamg/pyamg>. Release 4.0.
- 181 [8] B. SHRIMALI, V. LEFÈVRE, AND O. LOPEZ-PAMIES, *A simple explicit homog-  
 182 enization solution for the macroscopic elastic response of isotropic porous elas-  
 183 tomers*, Journal of the Mechanics and Physics of Solids, 122 (2019), pp. 364–380.
- 184 [9] L. TIAN, L. TEVET-DERE, K. BHATTACHARYA, ET AL., *Dielectric elastomer  
 185 composites*, Journal of the Mechanics and Physics of Solids, 60 (2012), pp. 181–  
 186 198.
- 187 [10] J. A. WEISS, B. N. MAKER, AND S. GOVINDJEE, *Finite element implemen-  
 188 tation of incompressible, transversely isotropic hyperelasticity*, Computer methods  
 189 in applied mechanics and engineering, 135 (1996), pp. 107–128.

190 **Appendix A. FE Solver.** The FE solver used to generate the numerical results  
 191 presented in section 3 uses Algorithm A.1 and the corresponding optimized parameters  
 192 used in the calculations are summarized in Table A.1.

---

**Algorithm A.1** FE Solver: Newton Multigrid

---

```

1: Define  $W := W(I_1, I_2, I_3)$  and calculate its derivatives wrt invariants  $\partial W/\partial I_i$ 
2: Define Element type , Quadrature rule, import geometry, mesh, connectivity ( $N_{el}$ )
   and assign boundary conditions
3: Create arrays for  $\mathbf{S}$ ,  $\mathbf{F}$ , dofstore
4: Initialize dofs (disp) and deformation gradient  $\mathbf{F}$ , Jacobian  $\partial^2 W/\partial \mathbf{F}^2$ , and residuals  $\partial W/\partial \mathbf{F}$ .
5: Define load steps nSteps and initialize step counter  $i$  to 0,  $i = \{1, 2, \dots, nSteps\}$ .

6: while  $i \leq nSteps$  do
7:   Initialize the Newton iteration counter iterN to 0, and apply disp= d
8:   Assemble the global Jacobian ( $\mathbf{K}$ ) and residual ( $\mathbf{r}$ ) as explained below
9:   while  $\|\mathbf{r}\{\mathbf{d}\}\|_2 > tolNR$  and iterNR  $\leq maxiters$  do
10:    Initialize element counter  $k_{el}$  to 0
11:    while  $k_{el} \leq N_{el}$  do
12:      At each gauss point inside the element  $k_{el}$  calculate  $\mathbf{F}$  using current value
         of  $\mathbf{d}$ 
13:      Compute  $\mathbf{S}$  from (1.2) at each quadrature point and append to the element
         tensor
14:      Integrate the Jacobian  $\partial^2 W/\partial \mathbf{F}^2$  and the residual  $\partial W/\partial \mathbf{F}$  over element
          $k_{el}$  and local quadrature points  $\{\xi, \eta\} \in (-1, 1)$  and calculate the element
         level tangent (for Newton)
15:      Assemble the element level energy  $W(\mathbf{F})$  at quadrature points and store
          $\mathbf{S}$  and  $\mathbf{F}$  for visualization
16:    end while
17:    Append the time taken for the assembly (timeAssemb) to global list
18:    Solve the linearized system  $\mathbf{K} \Delta \mathbf{d} = \mathbf{r}$  (1.9) using MG (aSA, RS, SA)
19:    Store the MG-residuals res and determine the average convergence factor  $\rho$ 
20:    Append the solve time tN to the global list timeStep
21:     $\mathbf{d} \leftarrow \mathbf{d} + \Delta \mathbf{d}$ 
22:    Compute the residual  $\mathbf{r}(\mathbf{d})$ 
23:    iterNR  $\leftarrow$  iterNR+1
24:  end while
25:   $i \leftarrow i + 1$ 
26: end while

```

---

193 Throughout the calculations we assume that the FE solution  $\tilde{\mathbf{u}}(\mathbf{X})$  is in the energy  
 194 space  $\mathcal{E}(\Omega)$  where

195 (A.1) 
$$\mathcal{E}(\Omega) := \mathbf{u} \ni \|\mathbf{F}(\mathbf{u})\| < +\infty \quad \text{and} \quad \int_{\Omega} W(\mathbf{F}) d\mathbf{X} < +\infty$$
  
 196

197 This allows to determine the convergence of the FE solution in the energy norm. Since  
 198 the problem (1.3) in general does not admit exact solution, we approximate the exact  
 199 solution by the FE computed on the finest mesh ( $\tilde{\Omega}$ ). The relative error in the energy  
 200 norm of the FE solution is given by

201 (A.2) 
$$e_{\mathcal{E}(\Omega)} = \sqrt{\frac{\mathcal{U} - \mathcal{U}_h}{\mathcal{U}}} \quad \text{where} \quad \mathcal{U} = \int_{\tilde{\Omega}} W(\mathbf{F}) d\mathbf{X}$$
  
 202

Solver	Optimized Parameters
aSA	initial_candidates = Bvec, num_candidates = 5, improvement_iters = 5, aggregate = 'naive'
GS	strength = 'evolution', CF = 'RS', max_levels = 50, pre/post-smoother = ('GS', 'symmetric'), max_coarse = 50
SA	smooth = ('energy'), strength = 'evolution', max_coarse = 50, max_levels= 15

Table A.1: Optimized Parameters for Multigrid Solvers used in [Algorithm A.1](#)